



Heard


A one-tap speech relay for dysarthria that learns your voice as you use it.

🏆 2nd place, Consumer Health Hackathon · London AI × Encode Club, backed by Juno (YC) & Anthropic · July 2026

They know exactly what they want to say. Their muscles won't execute it.

People with dysarthria build the sentence correctly. Friends and family understand them fine, because familiar listeners adapt. The barista doesn't. Neither does the GP receptionist, the taxi driver, or the person on the phone. Every one of those is a small failure, several times a day.

The problem isn't communication in general. **It's being understood by a stranger, once.**



**Two failed attempts at a coffee
counter.**

One tap. Coffee handed over.

One button. **Heard** does the rest.

Heard listens to the conversation in the background, recovers what the user meant, and speaks it aloud in a clean voice: matched to their age, gender and accent, or their own cloned voice where a pre-onset recording exists.

STEP 1

Let the other person speak first

Heard listens and shows their words on screen, building live conversation context.

STEP 2

Tap when you want to reply

Speak naturally. One big Speak button isolates your speech, then Heard shows its best recovery of what you meant.

STEP 3

Confirm, and Heard speaks for you

Tap Correct when it looks right. Heard plays your reply aloud in a clean voice.

It never speaks a confident error.

Trust is the product.

- When Heard is **confident**, it speaks straight away. The interaction stays at one tap.
- When it's **not sure, it doesn't guess**: it offers a couple of tappable options, plus a type-it fallback.
- Speaking a wrong sentence aloud in someone's own voice is worse than a two-second tap. The confidence gate is what makes users willing to hand Heard their voice at a counter, with a stranger watching.

Every correction is a training pair. The friction disappears with use.



- **Near-zero cold start.** Useful on day one via few-shot examples seeded from the TORGO dysarthric-speech dataset. No enrolment, no hours of recordings.
- **A word it has to ask about once, it simply speaks the next time.** The user's own confirmed pairs take over within a handful of uses.
- **Personal vocabulary compounds.** Per-user data beats any generic recogniser on that user, and the advantage grows with every conversation. The data is the moat.

Recovery, not recognition

Background ASR keeps live conversation context while the other person speaks.

Few-shot retrieval: the user's own confirmed pairs first, TORGO seed pool by text similarity, demographics as a tiny tie-breaker.

Confidence gate: confident → speak; unsure → tappable confirm cards.

ElevenLabs Scribe v2 transcribes the user's audio only, with word-level confidence.

LLM correction (Gemini 2.5 Flash) repairs the sentence using the retrieved pairs and conversation context.

ElevenLabs TTS speaks in a clean or cloned voice, and the confirmed pair is written back to the profile.

Parkinson's first, by design

Parkinsonian dysarthria is characteristic and consistent, so there is real structure for the system to learn. The same relay then serves dysarthria from cerebral palsy, MS, ALS/MND, stroke and TBI.

10M+



people live with Parkinson's
worldwide

~90%



of them are estimated to develop
speech impairment

Day 1



utility, with no enrolment or voice
banking required

Two camps, and neither is aimed at our moment

Personalised recognisers

Voiceitt, Google Project Euphonia

Make you train a model on hours of recordings before they're useful, then speak their single best guess.

Voice banking & AAC

Apple Personal Voice, ElevenLabs Impact Program

Give a voice to people who have lost speech entirely. They solve what comes out, not what the person who still speaks meant.

Heard

The only product for the person who **still speaks but isn't understood by a stranger**. Complementary to the voice players: a banked or cloned voice becomes the voice Heard speaks in.

Say it once. **Be understood everywhere.**

Once Heard has recovered what you meant, being understood doesn't have to stop at the person in front of you.



Live

The core loop, live today: the counter, the GP desk, the taxi, the phone call.



Share

Recovered messages sent as voice notes, text, a personalised avatar, or a generated video clip, straight to the channels the user already lives in.



My Words

The personal vocabulary, visible and growing: every confirmed correction the user has taught Heard.

Voice cloning and avatar/video use the user's own voice and face, with consent, as an assistive aid. Not impersonation of anyone else.

Built in a weekend. Working today.

- **Working product, live now** at heard-iota.vercel.app/live, with the learning loop wired end to end: listen, Speak, recover, confirm, speak aloud.
- 🥈 **2nd place** at the Consumer Health Hackathon (London AI × Encode Club, backed by Juno (YC) and Anthropic), London, July 2026. Judges singled out the UI.
- **Lean, proven stack:** Next.js on Vercel, Supabase, ElevenLabs Scribe v2 and TTS, LLM correction via Gemini 2.5 Flash. ASR and LLM keys never touch the browser.
- **Clear safety posture:** an assistive communication aid, not a medical device. No diagnosis, no clinical claims.

The ask: prototype to **Parkinson's** pilot

Raising a pre-seed round to harden the core loop, run a Parkinson's pilot with speech-and-language-therapy partners, and launch the consumer subscription. Clinic and SLT partnerships follow as the B2B channel.

Emre Yavuz

linkedin.com/in/dr-emre-yavuz-449216187

Sevara Bakhodirova

linkedin.com/in/sevara-bakhodirova-485102205

Rashad Hosseini

linkedin.com/in/rashad-hosseini

Maria Papageorgiou

linkedin.com/in/mariapapz