



2nd place, Consumer Health
Hackathon

London AI × Encode Club, backed by
Juno (YC) & Anthropic. July 2026.

A one-tap speech relay for dysarthria that **learns your voice** as you use it.

People with dysarthria know exactly what they want to say. Their muscles won't execute it. Friends and family adapt; the barista, the GP receptionist and the taxi driver do not. The problem isn't communication in general. **It's being understood by a stranger, once.**

WHAT HEARD DOES

Heard listens to the conversation in the background. When it's the user's turn, one big **Speak** button isolates their speech, recovers what they meant, and on a single confirm speaks it aloud in a clean voice: matched to their age, gender and accent by default, or their own cloned voice where a pre-onset recording exists.



When it's confident, it speaks straight away. When it's not sure, it never guesses: it offers tappable options plus a type-it fallback. **Trust is the product.**

THE MOAT: A LEARNING LOOP, NOT A MODEL

- **Near-zero cold start.** Useful on day one via few-shot examples seeded from the TORGO dysarthric-speech dataset. No enrolment, no hours of training recordings.
- **Every correction is a training pair.** A word Heard has to ask about once, it simply speaks the next time. The user's own confirmed pairs take over within a handful of uses.
- **Personal vocabulary compounds.** Per-user data beats any generic recogniser on that user, and the advantage grows with every conversation.

THE DEMO IN ONE LINE

Two failed attempts at a coffee counter. One tap. Coffee handed over.

MARKET

Parkinson's-first: its dysarthria is characteristic and consistent, so there is real structure to learn. The same relay serves cerebral palsy, MS, ALS/MND, stroke and TBI.

10M+

people live with Parkinson's worldwide

~90%

of them are estimated to develop speech impairment

Day 1

utility, no enrolment or voice banking required

WHY NOBODY ELSE OWNS THIS MOMENT

Personalised recognisers

Voiceitt · Euphonia

Hours of training recordings before they're useful; then they speak their single best guess.

Voice banking & AAC

Personal Voice · ElevenLabs

A voice for people who have lost speech entirely: what comes out, not what was meant.

Heard

Built for the person who still speaks but isn't understood by a stranger. Complementary: a banked voice becomes the voice Heard speaks in.

TRACTION & STATUS

- **Live today:** heard-iota.vercel.app/live, learning loop wired end to end.
- **2nd place** at the Consumer Health Hackathon (London AI × Encode Club), July 2026. Judges singled out the UI.
- **Stack:** Next.js on Vercel, Supabase, ElevenLabs Scribe v2 + TTS, Gemini 2.5 Flash correction. Keys server-side.

THE ASK

Raising a pre-seed round to take Heard from winning prototype to a Parkinson's pilot: consumer subscription first, with speech-therapy and clinic partnerships as the B2B channel. Roadmap adds sharing as voice notes, text, avatar and video, always with the user's own consented voice and likeness.

Team: Emre Yavuz · Sevara Bakhodirova · Rashad Hosseini · Maria Papageorgiou

An assistive communication aid, not a medical device: no diagnosis, no clinical claims.

emre@talktalk.net

github.com/eyavuz21/Heard · heard-iota.vercel.app/live